

How to migrate to AWS without a mental breakdown

Michal Artazov



Outline

- About signageOS
- How we migrated to AWS
- What went wrong
- How much did it cost
- What did we learn from it
- Some interesting ways we use AWS

signageOS

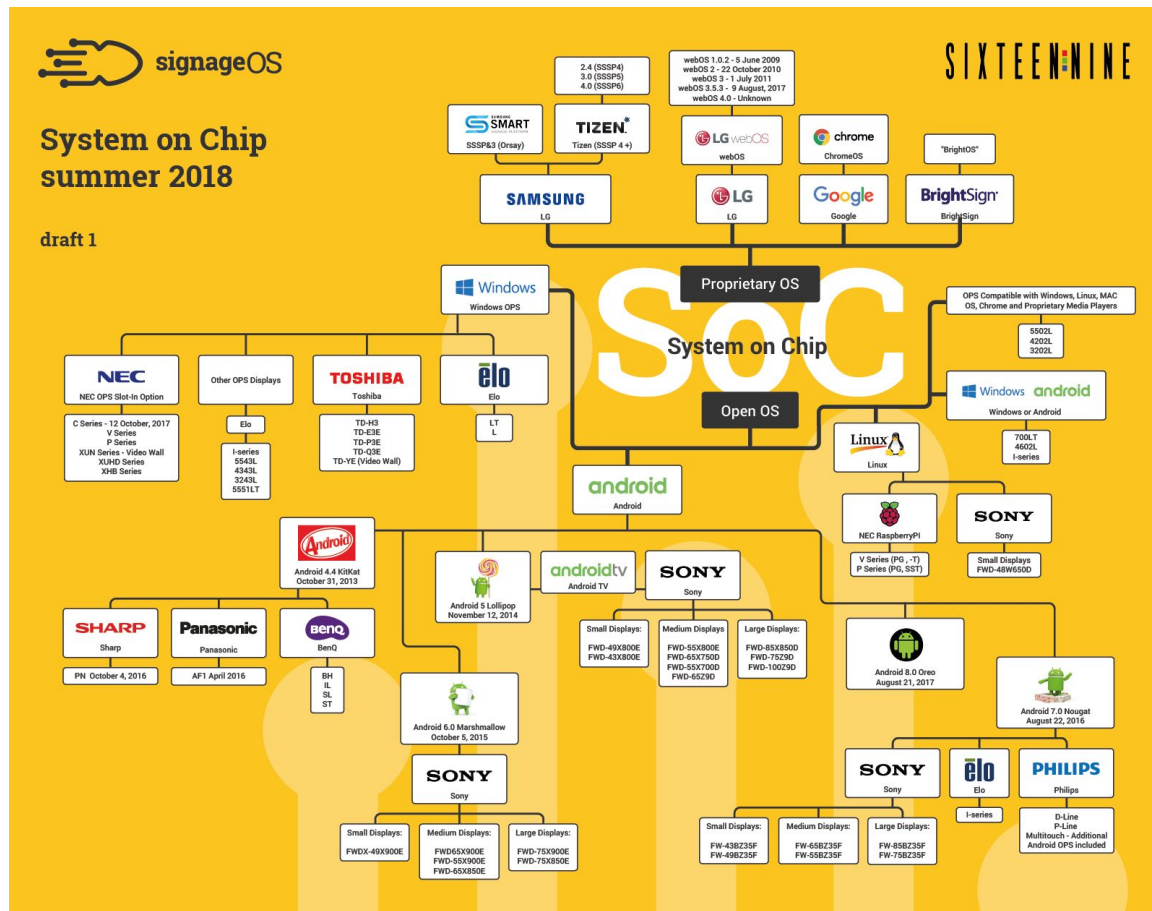
Out-of-box solution for device management, monitoring and content delivery designed for **digital signage**




What is digital signage



What is digital signage



What is signageOS



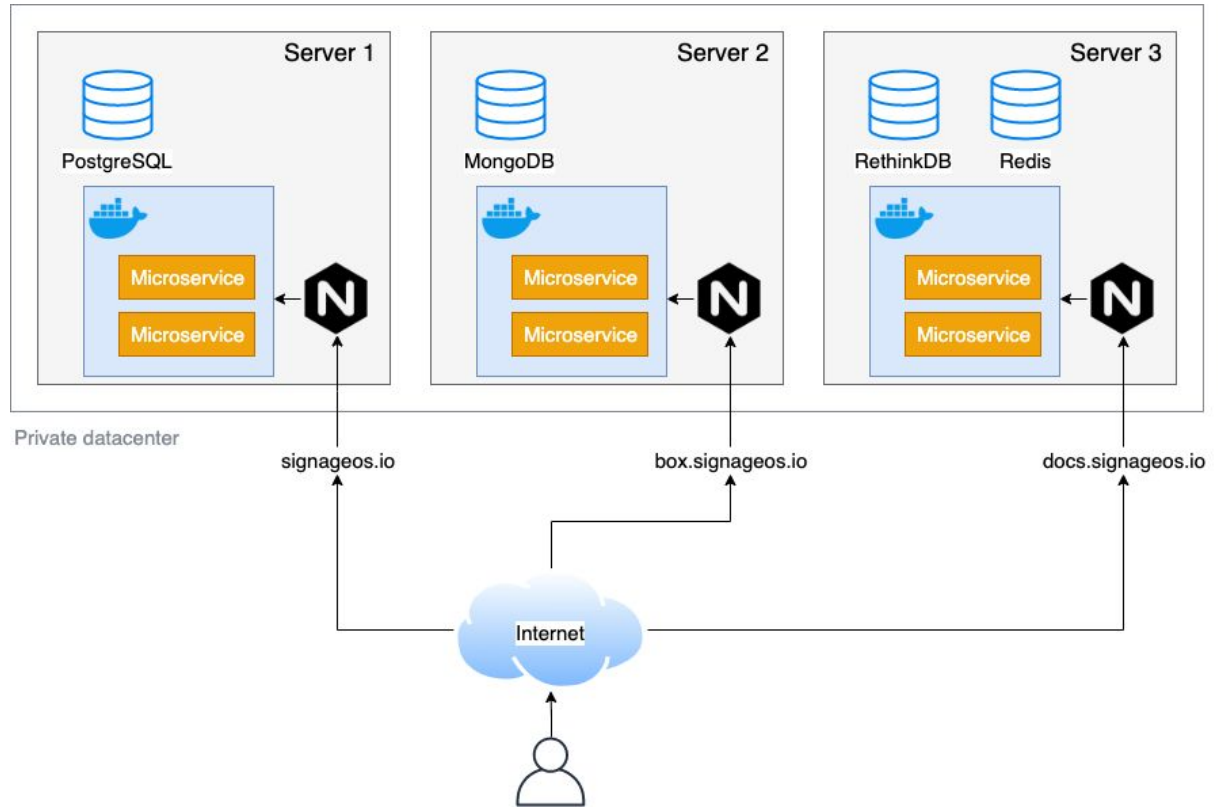
What is signageOS

- SDK + Cloud platform
- Install signageOS software on a device
- Write HTML/JS application using signageOS SDK
- Push your application to the cloud
- Application runs on the device in an embedded web browser
 - Show content
 - Control device
- Device can be controlled from the cloud
- Device constantly sends monitoring data to the cloud

What is signageOS

- Right now about 5000 devices in production and growing
- Customers all over the world (USA, Europe, Asia, Australia)

Before AWS



Before AWS

Pros

- Easy to understand

Cons

- Manual
- Not flexible
- Unable to scale

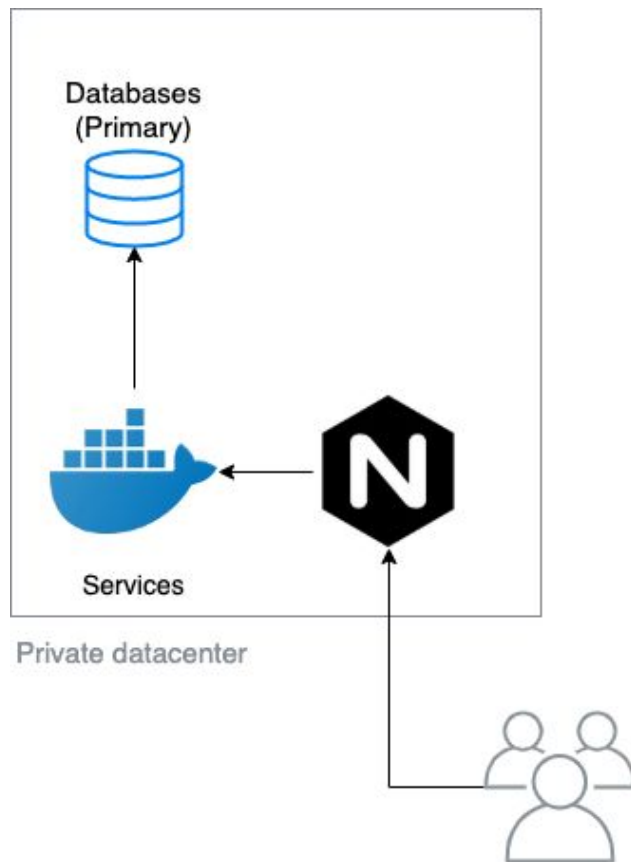
Why AWS

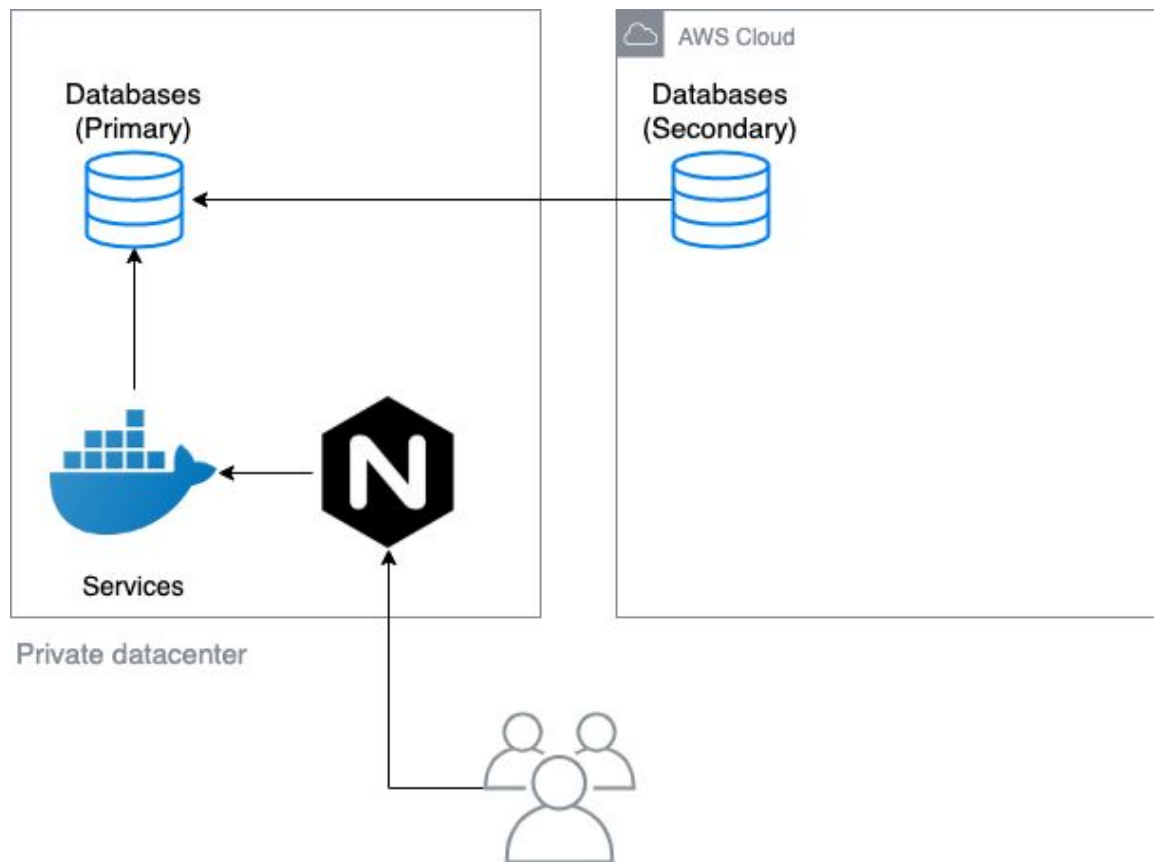
- Flexible
- Automation
- Better cost control
- More reliable

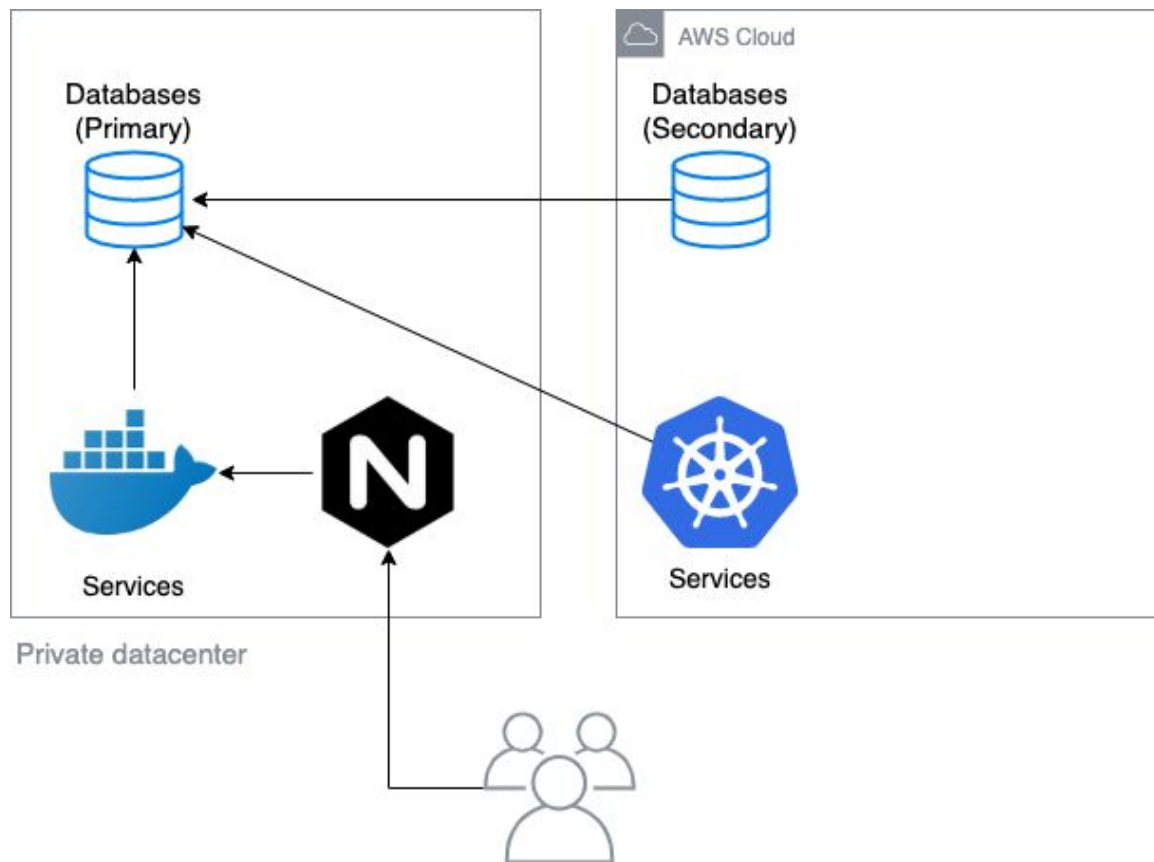
Migration - main objectives

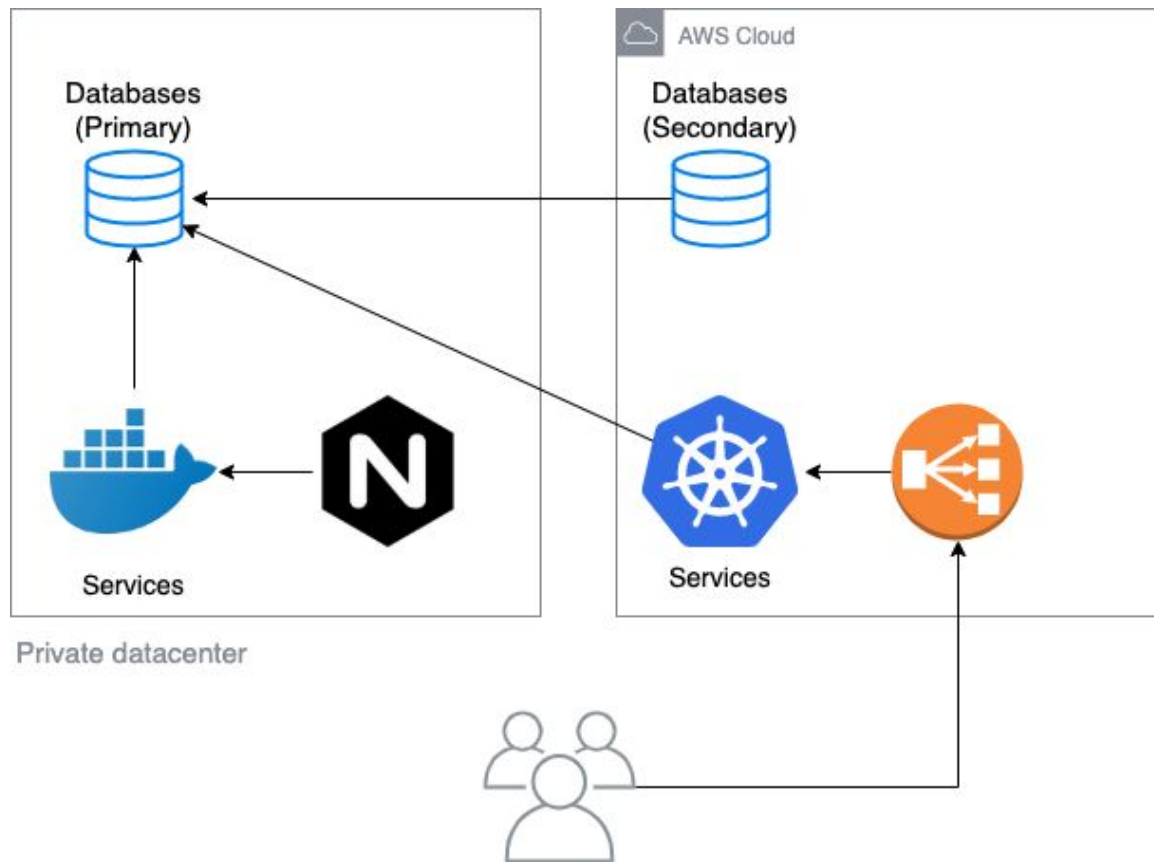
1. Replicate databases
2. Setup clone of the entire stack on AWS
3. Switch DNS to AWS
4. Cut off old database replicas
5. Profit












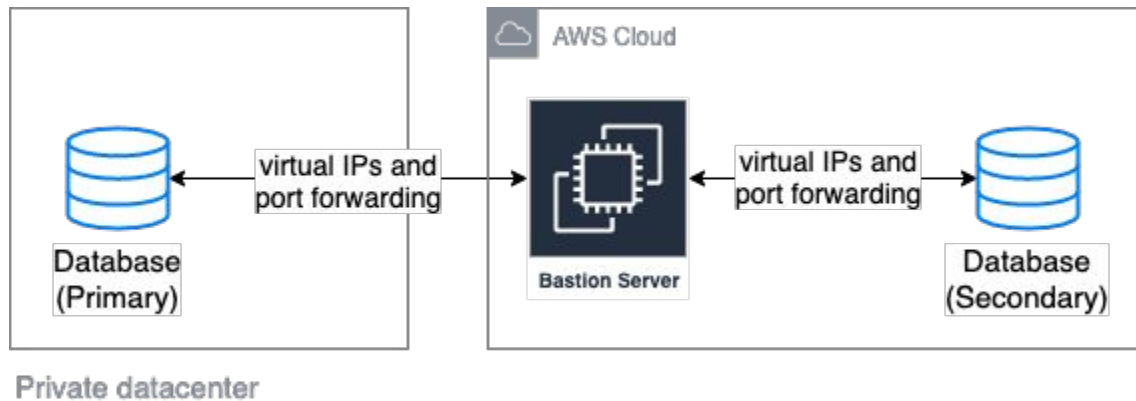








Database replication



Why not give an Elastic IP to every EC2 instance?

Too many instances

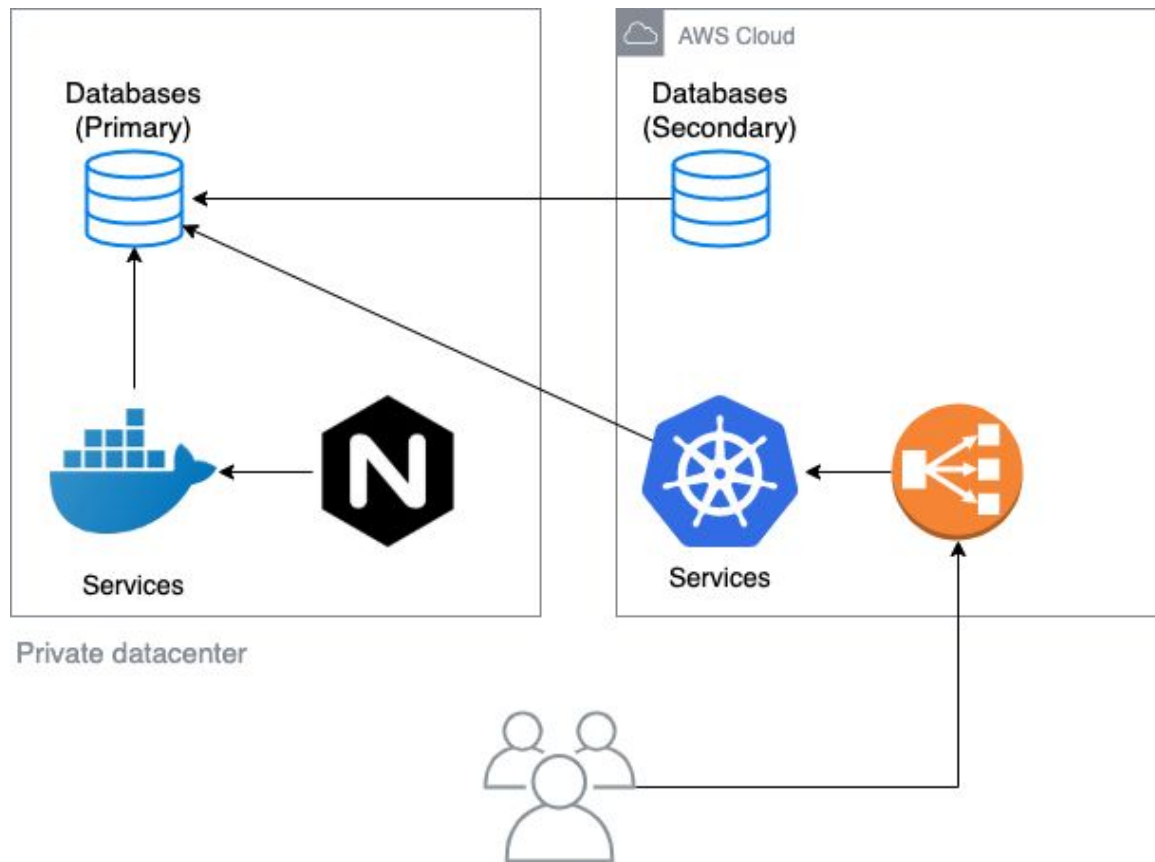
Kubernetes


- **EKS** (1.11)
- Already had Docker images for all services
- Used **Helm** to easier deploy to k8s
- Expose endpoints to the internet → [nginx ingress](#)
 - Creates ELB
 - Runs 1 or more instances of Nginx in k8s
 - Can run as DaemonSet → one instance per node

Helm

- In CI (Gitlab) build a Helm chart
- Publish chart to private repository
 - S3 - [Helm S3 plugin](#)
- Stack repo that combines all charts into **one big chart signageOS prod**
 - Defines all specific versions of services
 - Contains **complete production configuration**, encrypted with [git-secret](#)
 - Deploy changes to the production → push changes to Stack repo and deploy new chart







Switch databases primaries to AWS



Switch databases primaries to AWS

- Can't switch until all data is replicated
- Replication over the internet isn't great
- New data coming in faster than replicating
- EC2 instances running out of resources
 - Adding more resources → higher cost
- Result
 - Never ending cycle of death
 - Databases dropping connections
 - Core services fail
 - Pretty much whole system down

Solution


Sometimes the smart way isn't the best way


1. Only replicate the low volume collections
2. Temporarily stop all services that generate high volumes of data
 - Mostly read-only device monitoring and logs
 - New data not lost, just enqueued in RabbitMQ (millions of messages)
3. Rename high volume collections
4. Create new empty collections
5. Database finally replicated → switch primary to AWS
6. Manually copy the old data later











Final AWS structure



Costs

- \$3,000/month before AWS
- \$15,500 3 months of migration
- \$5,000 first month on AWS
- \$3,000/month now

**Other interesting facts about our
infrastructure**

Kubernetes nodes autoscaling

- Generic auto-scaling group
- RabbitMQ “dedicated” nodes
 - Run RabbitMQ instances (16 instances total)
 - Most of the time underutilized so kill other nodes and run as many services as possible here
 - When heavy traffic, start more generic nodes and move services there so these become dedicated to RabbitMQ
- Scale with CloudWatch
 - Collect metrics from nodes with [CloudWatch agent](#)
 - If both average CPU and memory consumption < 40% → less generic nodes
 - If average CPU or memory > 70% → more generic nodes

Serverless

- Lambda for user defined automatic tests
- Our custom simple testing framework that allows users to write automatic tests that run on device and validate the content
- Javascript code, interacts with the device through signageOS REST API to test the application that's running on the device
- Every test is converted into a lambda
 - Invoked from our services
 - Returns JSON with test results → save into our database
 - Test results available via REST API or management console (UI)



Sources

- <https://www.signageos.io>
- <https://github.com/nginxinc/kubernetes-ingress>
- <https://github.com/hypnoglow/helm-s3>
- <https://git-secret.io/>
- <https://docs.aws.amazon.com>

Contact

- Michal Artazov
- michal@signageos.io
- <https://www.signageos.io>

Questions